

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 from sklearn.cluster import KMeans
        5 import warnings
        6 df=pd.read_csv(r"C:\Users\saira\OneDrive\Desktop\heart.csv")
        7 df.head()
```

Out[1]:

	Age	Sex	ChestPain	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak	Slope	Ca	Thal
0	63	1	typical	145	233	1	2	150	0	2.3	3	0	fixed
1	67	1	asymptomatic	160	286	0	2	108	1	1.5	2	3	normal
2	67	1	asymptomatic	120	229	0	2	129	1	2.6	2	2	reversable
3	37	1	nonanginal	130	250	0	0	187	0	3.5	3	0	normal
4	41	0	nontypical	130	204	0	2	172	0	1.4	1	0	normal

```
In [2]: 1
        2 print(df.isnull().sum())
```

```
Age          0
Sex          0
ChestPain    0
RestBP       0
Chol         0
Fbs          0
RestECG      0
MaxHR        0
ExAng        0
Oldpeak      0
Slope        0
Ca           0
Thal         2
Target       0
dtype: int64
```

```
In [3]: 1 df=df.dropna()
```

In [4]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 301 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Age         301 non-null    int64
 1   Sex         301 non-null    int64
 2   ChestPain   301 non-null    object
 3   RestBP      301 non-null    int64
 4   Chol        301 non-null    int64
 5   Fbs         301 non-null    int64
 6   RestECG     301 non-null    int64
 7   MaxHR       301 non-null    int64
 8   ExAng       301 non-null    int64
 9   Oldpeak     301 non-null    float64
10   Slope       301 non-null    int64
11   Ca          301 non-null    int64
12   Thal        301 non-null    object
13   Target      301 non-null    int64
dtypes: float64(1), int64(11), object(2)
memory usage: 35.3+ KB
```

In [5]: 1 df.describe()

Out[5]:

	Age	Sex	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	C
count	301.000000	301.000000	301.000000	301.000000	301.000000	301.000000	301.000000	301.000000	301
mean	54.451827	0.681063	131.714286	246.936877	0.146179	0.990033	149.700997	0.325581	1
std	9.067258	0.466841	17.655729	51.859869	0.353874	0.994937	22.860817	0.469372	1
min	29.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0
25%	48.000000	0.000000	120.000000	211.000000	0.000000	0.000000	134.000000	0.000000	0
50%	56.000000	1.000000	130.000000	242.000000	0.000000	1.000000	153.000000	0.000000	0
75%	61.000000	1.000000	140.000000	275.000000	0.000000	2.000000	166.000000	1.000000	1
max	77.000000	1.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6

```
In [6]: 1 import pandas_profiling
        2 pandas_profiling.ProfileReport(df)
```

```
C:\Users\saira\AppData\Local\Temp\ipykernel_4792\1118840070.py:1: DeprecationWarning:
`import pandas_profiling` is going to be deprecated by April 1st. Please use `import y
data_profiling` instead.
  import pandas_profiling
```

Summarize dataset: 48/48 [00:06<00:00, 4.04it/s,
100% Completed]

Generate report structure: 100% 1/1 [00:06<00:00, 6.90s/it]

Render HTML: 100% 1/1 [00:01<00:00, 1.34s/it]

Overview

Dataset statistics

Number of variables	14
Number of observations	301
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	35.3 KiB
Average record size in memory	120.0 B

Variable types

Numeric	5
Categorical	9

Alerts

ChestPain is highly overall correlated with Target	High correlation
Target is highly overall correlated with ChestPain and 1 other fields (ChestPain, Thal)	High correlation
Thal is highly overall correlated with Target	High correlation

Out[6]:

```
In [7]: 1 d=df['Target'].value_counts()
        2 print(d)
```

```
0    163
1    138
Name: Target, dtype: int64
```

```
In [8]: 1 import seaborn as sns
        2 def plotTarget():
        3     sns.countplot(x='Target', data=df, ax=ax)
        4     for i, p in enumerate(ax.patches):
        5         count=df['Target'].value_counts().values[i]
        6         x=p.get_x()+ p.get_width() /2.
        7         y=p.get_height() + 3
        8         label='{:.2f}'.format(count / float(df.shape[0]))
        9         ax.text(x, y,label, ha='center')
       10
       11 fig_target,ax=plt.subplots(nrows=1, ncols=1, figsize=(5, 2))
       12 plotTarget()
```

```
In [9]: 1 df.corr()
```

C:\Users\saira\AppData\Local\Temp\ipykernel_4792\1134722465.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
df.corr()
```

Out[9]:

	Age	Sex	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak	
Age	1.000000	-0.098138	0.284734	0.208287	0.121670	0.149037	-0.395982	0.092985	0.203604	0
Sex	-0.098138	1.000000	-0.065284	-0.202126	0.041025	0.029016	-0.057065	0.140802	0.098482	0
RestBP	0.284734	-0.065284	1.000000	0.129371	0.178498	0.147089	-0.046402	0.065564	0.188801	0
Chol	0.208287	-0.202126	0.129371	1.000000	0.015762	0.171185	-0.005690	0.064250	0.044836	-0
Fbs	0.121670	0.041025	0.178498	0.015762	1.000000	0.079892	-0.012297	0.013534	0.004855	0
RestECG	0.149037	0.029016	0.147089	0.171185	0.079892	1.000000	-0.077950	0.092626	0.117580	0
MaxHR	-0.395982	-0.057065	-0.046402	-0.005690	-0.012297	-0.077950	1.000000	-0.386043	-0.349391	-0
ExAng	0.092985	0.140802	0.065564	0.064250	0.013534	0.092626	-0.386043	1.000000	0.287926	0
Oldpeak	0.203604	0.098482	0.188801	0.044836	0.004855	0.117580	-0.349391	0.287926	1.000000	0
Slope	0.162228	0.031571	0.117437	-0.004228	0.054079	0.140144	-0.393527	0.254076	0.576795	1
Ca	0.331939	0.100345	0.100535	0.106125	0.164689	0.127330	-0.256365	0.151731	0.274451	0
Target	0.224394	0.272006	0.151471	0.086762	0.015613	0.177049	-0.425870	0.427860	0.423894	0

```
In [10]: 1 x=df['ChestPain']
        2 x.value_counts()
```

```
Out[10]: asymptomatic    143
nonanginal      85
nontypical      50
typical         23
Name: ChestPain, dtype: int64
```

```
In [11]: 1 x=df['Thal']
          2 x.value_counts()
```

```
Out[11]: normal      166
          reversable  117
          fixed       18
          Name: Thal, dtype: int64
```

```
In [12]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 301 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Age        301 non-null   int64
 1   Sex        301 non-null   int64
 2   ChestPain  301 non-null   object
 3   RestBP     301 non-null   int64
 4   Chol       301 non-null   int64
 5   Fbs        301 non-null   int64
 6   RestECG    301 non-null   int64
 7   MaxHR      301 non-null   int64
 8   ExAng      301 non-null   int64
 9   Oldpeak    301 non-null   float64
10   Slope      301 non-null   int64
11   Ca         301 non-null   int64
12   Thal       301 non-null   object
13   Target     301 non-null   int64
dtypes: float64(1), int64(11), object(2)
memory usage: 43.4+ KB
```

```
In [13]: 1 df_X= df.loc[:, df.columns != 'Target']
          2 df_y= df.loc[:, df.columns == 'Target']
```

```
In [14]: 1 from sklearn.compose import make_column_selector as selector
          2
          3 numerical_columns_selector = selector(dtype_exclude=object)
          4 categorical_columns_selector = selector(dtype_include=object)
          5
          6 numerical_columns = numerical_columns_selector(df_X)
          7 categorical_columns = categorical_columns_selector(df_X)
```

```
In [18]: 1 from sklearn.preprocessing import LabelEncoder, StandardScaler
```

```
In [26]: 1 categorical_preprocessor = Pipeline(steps=[('label_encoder', LabelEncoder())])
          2 numerical_preprocessor = StandardScaler()
```

```
In [27]: 1 categorical_columns,numerical_columns
```

```
Out[27]: (['ChestPain', 'Thal'],
          ['Age',
           'Sex',
           'RestBP',
           'Chol',
           'Fbs',
           'RestECG',
           'MaxHR',
           'ExAng',
           'Oldpeak',
           'Slope',
           'Ca'])
```

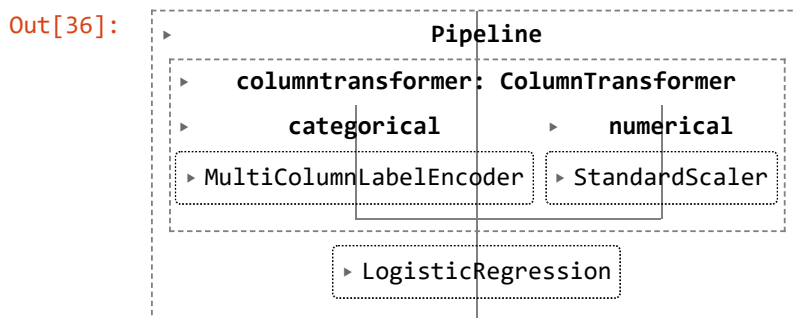
```
In [34]: 1 from sklearn.base import TransformerMixin
2 from sklearn.preprocessing import LabelEncoder
3 from sklearn.utils.validation import check_is_fitted
4 from sklearn.compose import ColumnTransformer
5 class MultiColumnLabelEncoder(TransformerMixin):
6     def __init__(self):
7         self.encoders = {}
8
9     def fit(self, X, y=None):
10        for col in X.columns:
11            self.encoders[col] = LabelEncoder().fit(X[col])
12        return self
13
14    def transform(self, X):
15        check_is_fitted(self, 'encoders')
16        output = X.copy()
17        for col, encoder in self.encoders.items():
18            output[col] = encoder.transform(X[col])
19        return output
20
21    def get_params(self, deep=True):
22        return {}
23
24    categorical_preprocessor = MultiColumnLabelEncoder()
25    numerical_preprocessor = StandardScaler()
26
27    preprocessor = ColumnTransformer(
28        transformers=[
29            ("categorical", categorical_preprocessor, categorical_columns),
30            ("numerical", numerical_preprocessor, numerical_columns),
31        ]
32    )
```

```
In [35]: 1 from sklearn.model_selection import train_test_split
2 X_train,X_test, y_train, y_test=train_test_split(df_X,df_y, test_size = 0.25, random_state=42)
3 columns = X_train.columns
```

```
In [36]: 1 from sklearn.pipeline import Pipeline
2 from sklearn.linear_model import LogisticRegression
3 pipeline = Pipeline(steps=[('columntransformer', preprocessor), ('LogisticRegression', LogisticRegression())])
4 pipeline.fit(X_train, y_train)
```

C:\Users\saira\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1184: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)



```
In [37]: 1 pre=pipeline.predict(X_test)
```

	precision	recall	f1-score	support
0	0.90	0.82	0.86	44
1	0.78	0.88	0.82	32
accuracy			0.84	76
macro avg	0.84	0.85	0.84	76
weighted avg	0.85	0.84	0.84	76

```
[[36  8]
 [ 4 28]]
```

[illegible]

```
1 from sklearn.metrics import classification_report, confusion_matrix
2 print(classification_report(y_pred, y_test))
3 print(confusion_matrix(y_pred, y_test))
```

	precision	recall	f1-score	support
0	0.82	0.85	0.84	39
1	0.83	0.81	0.82	37
accuracy			0.83	76
macro avg	0.83	0.83	0.83	76
weighted avg	0.83	0.83	0.83	76

```
[[33  6]
 [ 7 30]]
```

```

In [ ]: 1 """
2 from sklearn.model_selection import GridSearchCV
3 from sklearn.pipeline import Pipeline
4 import lightgbm as lgb
5
6 # Define the pipeline with hyperparameter tuning
7 param_grid = {
8     'lgbmclassifier__num_leaves': [31, 62, 127],
9     'lgbmclassifier__learning_rate': [0.2, 0.119, 0.21],
10    'lgbmclassifier__n_estimators': [91, 92],
11    'lgbmclassifier__max_depth': [3, 5, 7],
12    'lgbmclassifier__min_child_samples': [2, 5, 10],
13    'lgbmclassifier__min_child_weight': [0.001, 0.1, 1.0],
14    'lgbmclassifier__subsample': [0.5, 0.8, 1.0],
15    'lgbmclassifier__colsample_bytree': [0.5, 0.8, 1.0],
16    'lgbmclassifier__reg_alpha': [0.0, 0.1, 1.0],
17    'lgbmclassifier__reg_lambda': [0.0, 0.1, 1.0],
18    'lgbmclassifier__feature_fraction': [0.5, 0.8, 1.0],
19    'lgbmclassifier__bagging_fraction': [0.5, 0.8, 1.0],
20    'lgbmclassifier__bagging_freq': [1, 3, 5],
21    'lgbmclassifier__early_stopping_round': [5, 10, 20]
22 }
23
24 pipeline = Pipeline(steps=[('columntransformer', preprocessor),
25                             ('lgbmclassifier', lgb.LGBMClassifier())])
26
27 grid_search = GridSearchCV(pipeline, param_grid, cv=5, scoring='f1_macro')
28 grid_search.fit(X_train, y_train)
29
30 print("Best parameters:", grid_search.best_params_)
31 print("Best score:", grid_search.best_score_)

```

C:\Users\saira\anaconda3\Lib\site-packages\sklearn\preprocessing_label.py:97: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

C:\Users\saira\anaconda3\Lib\site-packages\sklearn\preprocessing_label.py:132: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)

C:\Users\saira\anaconda3\Lib\site-packages\sklearn\preprocessing_label.py:97: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

C:\Users\saira\anaconda3\Lib\site-packages\sklearn\preprocessing_label.py:132: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)

C:\Users\saira\anaconda3\Lib\site-packages\sklearn\preprocessing_label.py:97: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```

In [ ]: 1 y_pred = grid_search.predict(X_test)

```

```

In [ ]: 1 from sklearn.metrics import classification_report, confusion_matrix
2 print(classification_report(y_pred, y_test))
3 print(confusion_matrix(y_pred, y_test))

```


In []:

1